

Assignment 2

July 20, 2019

*You are currently looking at **version 1.0** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.*

1 Assignment 2 - Introduction to NLTK

In part 1 of this assignment you will use nltk to explore the Herman Melville novel Moby Dick. Then in part 2 you will create a spelling recommender function that uses nltk to find words similar to the misspelling.

1.1 Part 1 - Analyzing Moby Dick

```
In [13]: import nltk
import pandas as pd
import numpy as np
nltk.download('punkt')
nltk.download('averaged_perceptron_tagger')

# If you would like to work with the raw text you can use 'moby_raw'
with open('moby.txt', 'r') as f:
    moby_raw = f.read()

# If you would like to work with the novel in nltk.Text format you can use 'text1'
moby_tokens = nltk.word_tokenize(moby_raw)
text1 = nltk.Text(moby_tokens)

[nltk_data] Downloading package punkt to /home/jovyan/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data] /home/jovyan/nltk_data...
[nltk_data] Package averaged_perceptron_tagger is already up-to-
[nltk_data] date!
```

1.1.1 Example 1

How many tokens (words and punctuation symbols) are in text1?

This function should return an integer.

```
In [ ]: def example_one():  
  
    return len(nltk.word_tokenize(moby_raw)) # or alternatively len(text1)  
  
example_one()
```

1.1.2 Example 2

How many unique tokens (unique words and punctuation) does text1 have?

This function should return an integer.

```
In [ ]: def example_two():  
  
    return len(set(nltk.word_tokenize(moby_raw))) # or alternatively len(set(text1))  
  
example_two()
```

1.1.3 Example 3

After lemmatizing the verbs, how many unique tokens does text1 have?

This function should return an integer.

```
In [ ]: from nltk.stem import WordNetLemmatizer  
  
def example_three():  
  
    lemmatizer = WordNetLemmatizer()  
    lemmatized = [lemmatizer.lemmatize(w, 'v') for w in text1]  
  
    return len(set(lemmatized))  
  
example_three()
```

1.1.4 Question 1

What is the lexical diversity of the given text input? (i.e. ratio of unique tokens to the total number of tokens)

This function should return a float.

```
In [4]: def answer_one():  
  
    return (len(set(text1))/len(text1))  
  
answer_one()
```

```
Out[4]: 0.08139566804842562
```

1.1.5 Question 2

What percentage of tokens is 'whale' or 'Whale'?

This function should return a float.

```
In [17]: def answer_two():  
  
         whale=[w for w in text1 if (w=='whale') | (w=='Whale')]  
         return (len(whale)/len(text1))*100  
  
         answer_two()
```

```
Out[17]: 0.4125668166077752
```

1.1.6 Question 3

What are the 20 most frequently occurring (unique) tokens in the text? What is their frequency?

*This function should return a list of 20 tuples where each tuple is of the form (token, frequency).
The list should be sorted in descending order of frequency.*

```
In [22]: from nltk.probability import FreqDist  
         def answer_three():  
  
             dist=FreqDist(text1)  
  
             return list((sorted(dist.items(), key = lambda kv:(kv[1], kv[0]),reverse=True))[:20])  
  
         answer_three()
```

```
Out[22]: [(' ', 19204),  
          ('the', 13715),  
          ('.', 7308),  
          ('of', 6513),  
          ('and', 6010),  
          ('a', 4545),  
          ('to', 4515),  
          (';', 4173),  
          ('in', 3908),  
          ('that', 2978),  
          ('his', 2459),  
          ('it', 2196),  
          ('I', 2097),  
          ('!', 1767),  
          ('is', 1722),  
          ('--', 1713),  
          ('with', 1659),  
          ('he', 1658),  
          ('was', 1639),  
          ('as', 1620)]
```

1.1.7 Question 4

What tokens have a length of greater than 5 and frequency of more than 150?

This function should return an alphabetically sorted list of the tokens that match the above constraints. To sort your list, use `sorted()`

```
In [30]: from nltk.probability import FreqDist
def answer_four():

    dist=FreqDist(text1)
    return sorted(set([w for w in text1 if (len(w)>5) & (dist[w]>150)]))

answer_four()
```

```
Out[30]: ['Captain',
          'Pequod',
          'Queequeg',
          'Starbuck',
          'almost',
          'before',
          'himself',
          'little',
          'seemed',
          'should',
          'though',
          'through',
          'whales',
          'without']
```

1.1.8 Question 5

Find the longest word in text1 and that word's length.

This function should return a tuple (`longest_word`, `length`).

```
In [42]: def answer_five():

    maxlen=max([len(w) for w in text1 ])
    longw= [w for w in text1 if len(w)==maxlen]
    return (longw[0],maxlen)

answer_five()
```

```
Out[42]: ('twelve-o'clock-at-night', 23)
```

1.1.9 Question 6

What unique words have a frequency of more than 2000? What is their frequency?

"Hint: you may want to use `isalpha()` to check if the token is a word and not punctuation."

This function should return a list of tuples of the form (`frequency`, `word`) sorted in descending order of frequency.

```
In [56]: from nltk.probability import FreqDist
def answer_six():

    dist=FreqDist(text1)
    return sorted([(fre[1],fre[0]) for fre in dist.items() if (fre[1]>2000) and (fre[0]

answer_six()

Out[56]: [(13715, 'the'),
(6513, 'of'),
(6010, 'and'),
(4545, 'a'),
(4515, 'to'),
(3908, 'in'),
(2978, 'that'),
(2459, 'his'),
(2196, 'it'),
(2097, 'I')]
```

1.1.10 Question 7

What is the average number of tokens per sentence?

This function should return a float.

```
In [63]: from nltk.tokenize import sent_tokenize
from nltk.tokenize import word_tokenize
def answer_seven():

    moby_sentence = sent_tokenize(moby_raw)
    add=np.mean([len(word_tokenize(w)) for w in moby_sentence])
    return add

answer_seven()
```

```
Out[63]: 25.881952902963864
```

1.1.11 Question 8

What are the 5 most frequent parts of speech in this text? What is their frequency?

This function should return a list of tuples of the form (part_of_speech, frequency) sorted in descending order of frequency.

```
In [16]: from nltk.probability import FreqDist
import nltk
def answer_eight():

    pos=nltk.pos_tag(text1)
    dist=FreqDist([f for (p,f) in pos])
    count=dist.most_common()[5]
```

```

        return count

    answer_eight()

Out[16]: [('NN', 32730), ('IN', 28657), ('DT', 25867), ('', 19204), ('JJ', 17620)]

```

1.2 Part 2 - Spelling Recommender

For this part of the assignment you will create three different spelling recommenders, that each take a list of misspelled words and recommends a correctly spelled word for every word in the list.

For every misspelled word, the recommender should find the word in `correct_spellings` that has the shortest distance*, and starts with the same letter as the misspelled word, and return that word as a recommendation.

*Each of the three different recommenders will use a different distance measure (outlined below).

Each of the recommenders should provide recommendations for the three default words provided: `['cormulent', 'incendenece', 'validate']`.

```

In [4]: from nltk.corpus import words
        nltk.download('words')

        correct_spellings = words.words()

[nltk_data] Downloading package words to /home/jovyan/nltk_data...
[nltk_data] Unzipping corpora/words.zip.

```

1.2.1 Question 9

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Jaccard distance on the trigrams of the two words.

This function should return a list of length three: `['cormulent_reccommendation', 'incendenece_reccommendation', 'validate_reccommendation']`.

```

In [7]: def answer_nine(entries=['cormulent', 'incendenece', 'validate']):

        correct_spellings = words.words()
        jd=1
        corrected=[]
        for w in entries:
            first_let=lambda x: x if (x[0] == w[0]) else ''

            same_start=np.array(list(map(first_let,correct_spellings)))
            same_start=(same_start[same_start!=''])

            set1 = set(nltk.ngrams(w,n=3))
            compute_jd = lambda x: (nltk.distance.jaccard_distance(set1,set(nltk.ngrams(x,n=

```

```

jd=list(map(compute_jd,same_start))
corrected.append(sorted(jd)[0][1])

return corrected

```

```

answer_nine()

```

```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:13: DeprecationWarning: generator '
del sys.path[0]

```

```

Out[7]: ['corpulent', 'indecence', 'validate']

```

1.2.2 Question 10

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Jaccard distance on the 4-grams of the two words.

This function should return a list of length three: ['cormulent_reccomendation', 'incendenece_reccomendation', 'validrate_reccomendation'].

```

In [9]: def answer_ten(entries=['cormulent', 'incendenece', 'validrate']):

```

```

    correct_spellings = words.words()
    jd=1
    corrected=[]
    for w in entries:
        first_let=lambda x: x if (x[0] == w[0]) else ''

        same_start=np.array(list(map(first_let,correct_spellings)))
        same_start=(same_start[same_start!=''])

        set1 = set(nltk.ngrams(w,n=4))
        compute_jd = lambda x: (nltk.distance.jaccard_distance(set1,set(nltk.ngrams(x,n=
        jd=list(map(compute_jd,same_start))
        corrected.append(sorted(jd)[0][1])

    return corrected

```

```

answer_ten()

```

```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:13: DeprecationWarning: generator '
del sys.path[0]

```

```

Out[9]: ['cormus', 'incendiary', 'valid']

```

1.2.3 Question 11

For this recommender, your function should provide recommendations for the three default words provided above using the following distance metric:

Edit distance on the two words with transpositions.

This function should return a list of length three: ['cormulent_reccomendation', 'incendenece_reccomendation', 'validrate_reccomendation'].

```
In [10]: def answer_eleven(entries=['cormulent', 'incendenece', 'validrate']):
    correct_spellings = words.words()
    jd=1
    corrected=[]
    for w in entries:
        first_let=lambda x: x if (x[0] == w[0]) else ''

        same_start=np.array(list(map(first_let,correct_spellings)))
        same_start=(same_start[same_start!=''])

        compute_ed = lambda x: (nltk.distance.edit_distance(x,w),x)
        ed=list(map(compute_ed,same_start))
        corrected.append(sorted(ed)[0][1])

    return corrected

answer_eleven()

Out[10]: ['corpulent', 'intendence', 'validate']
```